

Taxonomy of Parameter-Efficient Fine-Tuning

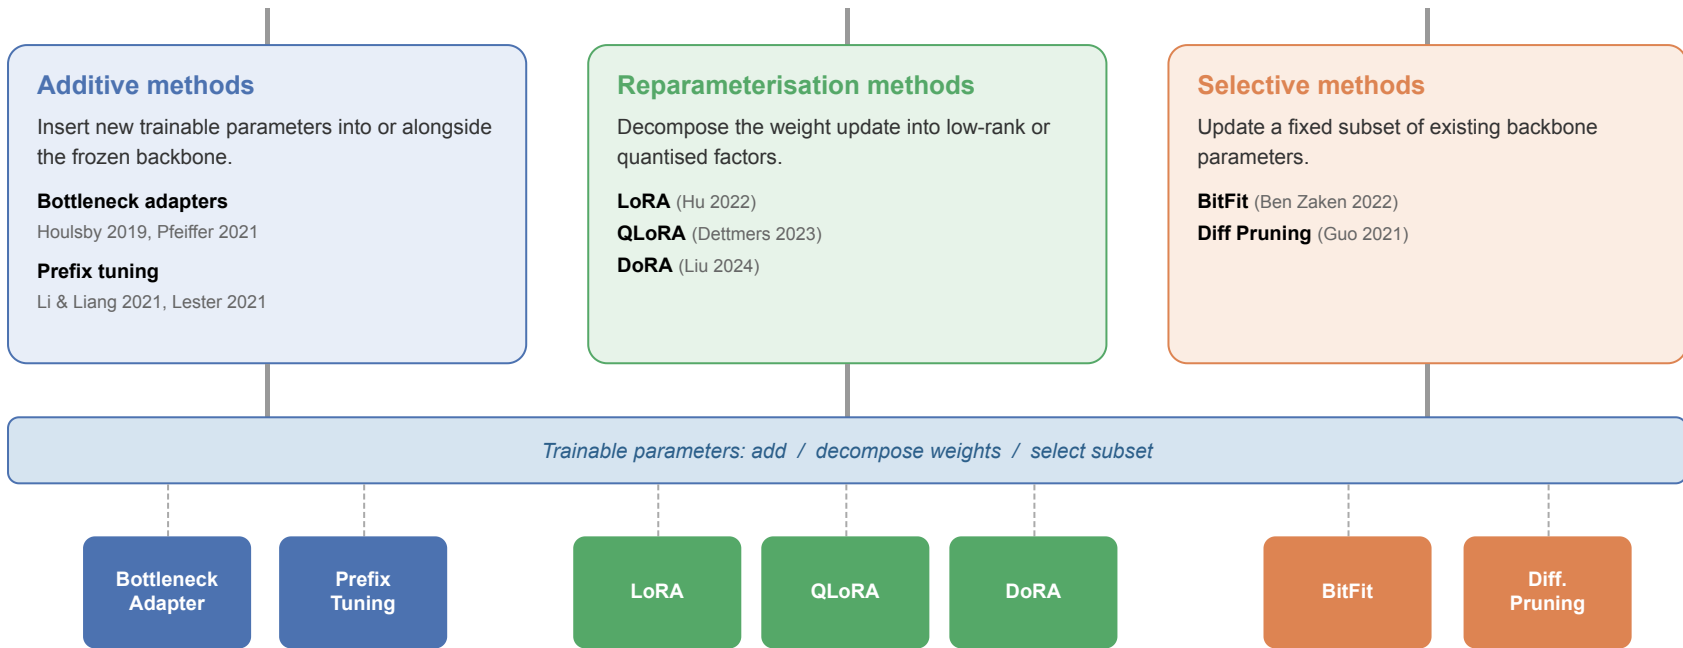


Figure adapted from the unified framework of Han et al. (2024)

Trade-off: additive methods add the most inference-time overhead but are the most modular (swappable per task); reparameterisation methods (LoRA family) add zero inference overhead after weight merging; selective methods add no new parameters but require access to update the backbone directly.